

The Per-Chunk Tax

A cross-ecosystem survey of HTTP/1.1 chunked-transfer-encoding amplification in production servers

Michael J. Bommarito II

June 2, 2026

Abstract

HTTP/1.1 chunked transfer encoding [1] is the standard mechanism by which a client transmits a request body whose length is unknown when the request line is sent. The specification defines the wire format but not the granularity at which a parser must deliver decoded body bytes to the application handler. We measure that delivery granularity, end to end and under controlled network conditions, across **27 production HTTP/1.1 servers** spanning **7 language ecosystems** (Python, Go, Rust, JVM, Node.js, BEAM, Ruby) and three C servers (nginx, Apache httpd, HAProxy).

We find that **exactly one HTTP/1.1 server** in our sample (Microsoft’s Kestrel on ASP.NET Core, via the `System.IO.Pipelines` substrate) *moves the per-chunk cost off the `async/await` scheduling boundary*: its `Http1ChunkedEncodingMessageBody.PumpAsync` drains the socket-readable buffer into a `Pipe`, loops through every chunk header present, then wakes the application handler with a single coalesced `PipeReader.ReadAsync` result. The HTTP/2 path of the same server (Wave 3) has no pump-equivalent and falls back to per-DATA-frame dispatch like every other HTTP/2 implementation we tested. Kestrel still pays a per-chunk parsing cost inside the H/1 pump (3.6 μ s/chunk steady-state), but eliminates the per-chunk task switch that dominates the other servers’ cost. Every other server we measure dispatches one parser callback per chunk and emits one application-receive event per chunk; this costs between 2.1 μ s (uvicorn + httptools) and 114 μ s (nginx as origin) of single-core time per chunk under our standardized probe (we define exactly what this number measures, and which cells are CPU- versus wall-derived, in Section 4.4).

A single superlinear mechanism, exhibited by the three Node `http.Server`-based servers (the built-in `http.Server` plus `express` and `fastify` atop it), drives genuine $\Theta(N^2)$ wall-time scaling (measured exponent ~ 1.7 – 2.0 in our N range) via the `process.nextTick` cost of the `Readable` substrate. A 1.5 MiB attacker request drives Node `http.Server` to 282 s of single-core wall time (Mode A, $N=250,000$, server pinned at one vCPU and saturated for the full run, so wall $\approx$ CPU); the `express` and `fastify` $N=250,000$ cells are extrapolated from a quadratic fit to the measured smaller cells, not measured. Elixir’s `bandit` is *not* quadratic: its measured exponent is ~ 1.05 and its per-chunk cost *falls* with N ($155 \rightarrow 148 \rightarrow 138$ μ s/chunk over $N=50k \rightarrow 100k \rightarrow 250k$ under Mode B). It is a high-constant linear case (~ 138 μ s/chunk via an `erlang:iolist_size/iolist` walk per chunk in `do_read_chunk!/4`), not a $\Theta(N^2)$ blow-up.

The “deploy behind nginx” mitigation widely advised in framework documentation is empirically confirmed. With default `proxy_request_buffering` on, an upstream-instrumented end-to-end probe (prober $\rightarrow$ nginx-proxy $\rightarrow$ instrumented Python upstream) shows the upstream receives *exactly one* `Content-Length`-framed request with no `Transfer-Encoding` header and *one* `recv()` covering both headers and the entire body. With `proxy_request_buffering` off the same probe forwards the chunked stream unchanged, with the upstream seeing multiple `recv()`s and a `Transfer-Encoding: chunked` header. Apache `mod_proxy_http` is documented to behave the same way (not directly measured). HAProxy does *not* aggregate: it streams by default, and `http-buffer-request` only waits for one buffer’s worth, not a per-chunk accumulator. Direct-exposed servers, HAProxy-fronted deployments, and any deployment with buffering disabled inherit the full per-chunk cost.

A Wave 3 measurement set extends the survey to HTTP/2 DATA frames and WebSocket text frames. Every HTTP/2 server in our sample (6 servers including Kestrel-h2, hypercorn-h2, node-h2, vertx-h2, go-h2c, rust-hyper-h2) exhibits the same per-frame amplification shape as HTTP/1 chunked-TE, at 4–10 μ s/frame; HTTP/2’s flow control bounds bytes but not frames. WebSocket is the only protocol

in which most ecosystems already batch correctly (node-ws, gorilla, rust-tungstenite, Kestrel-WS at 0.1–0.7 μ s/frame); Python ASGI’s WebSocket implementations (uvicorn + websockets / wsproto) are the at-risk outlier.

We characterize the root cause as a parser/dispatcher boundary design choice rather than a spec violation, taxonomize the mitigations available today, and report the responses of upstream maintainers including the closure of hyper issue #4008 [2] ("Provide built-in chunked request limits and CPU-safe streaming helpers") as `not_planned` with the labelled state `S-waiting-on-author`. We argue that opt-in aggregation primitives modeled on Kestrel’s Pipelines pump should be exposed by every event-loop HTTP server.

1 Introduction

The bug class. An HTTP/1.1 server that receives a request body via chunked transfer encoding must, somewhere on the data path, transition decoded body bytes from the wire-level parser into application code. The specification [1] fixes the wire format (`<hex-size>\r\n<bytes>\r\n`, terminated by a zero-length chunk) but is silent on how a parser should schedule callbacks into the application: per chunk, per `recv()` batch, per fixed-byte threshold, or fully buffered.

We measure which of those choices each production HTTP server has adopted as the default, and the cost in single-core CPU of the chosen granularity under a pathological-but-RFC-compliant input shape: an N -byte body encoded as N one-byte chunked-TE chunks.

Headline finding. Of 27 production HTTP/1.1 servers measured across 7 language ecosystems plus three C servers (a further 6 HTTP/2 and 6 WebSocket servers are measured in Wave 3, 39 records total):

- **1 HTTP/1.1 server**, Kestrel on ASP.NET Core, moves the per-chunk cost off the `async/await` scheduling boundary via `System.IO.Pipelines`. (Four of the six WebSocket servers also batch correctly; see Section 5.) Kestrel still pays a per-chunk parsing cost (3.6 μ s/chunk steady-state) but the resulting end-to-end cost is dominated by the synchronous parser loop rather than a per-chunk task switch.
- **23 servers** are linear-but-per-chunk vulnerable, with per-chunk costs spanning a factor of $\sim 65\times$ (2.1 μ s/chunk for `uvicorn+httptools` up to 138 μ s/chunk for the `bandit` iolist outlier; `nginx` as origin sits at 114 μ s/chunk).
- **3 servers** exhibit superlinear scaling consistent with an asymptotic $\Theta(N^2)$ mechanism observed in the profile, with measured exponents in the 1.7–2.0 range over our N range: Node’s built-in `http.Server`, `express`, and `fastify` – all three sharing the same `http.Server/1lhttp` Readable substrate. Elixir’s `bandit` is a high-constant *linear* case (measured exponent ~ 1.05 ; per-chunk cost *falls* from 155 to 138 μ s as N grows from 50k to 250k under Mode B) and is classified VULNERABLE-PER-CHUNK, not quadratic.

Figure 1 ranks all 27 HTTP/1.1 servers by per-chunk cost under the canonical Mode B probe (paced 100 μ s gap, $N=250,000$, 1 vCPU container).

Contributions.

1. The first systematic cross-ecosystem measurement of chunked-TE delivery granularity in production HTTP servers, spanning 7 language ecosystems and three C servers (Section 5, Section 6).
2. Identification of the superlinear route exhibited by the three Node `http.Server`-based servers (the Readable substrate’s `process.nextTick` drain), and a separate high-constant linear cost in `bandit`’s iolist accumulator (measured exponent ~ 1.05) that does *not* reach the $\Theta(N^2)$ regime in our N range.




Figure 1: Per-chunk server cost across all 27 HTTP/1.1 servers measured (Mode B: paced 100 μ s gap, N=250,000 one-byte chunks, 1 vCPU container). Log-scale x-axis. Spans nearly three orders of magnitude. Kestrel (BATCHES-CORRECTLY, leftmost) is the single existence proof among HTTP/1.1 servers.

3. Identification of an HTTP/1.1 existence proof (`System.IO.Pipelines/Kestrel`) that the bug class can be structurally closed at the server layer – corroborated by four WebSocket implementations (`node-ws`, `gorilla`, `tungstenite`, `Kestrel-WS`) that already batch by default; the design is not novel to this work but its security relevance has not been publicly named.
4. A characterization of the "deploy behind nginx" mitigation: default nginx *does* aggregate chunks for the upstream (Section 9), but HAProxy does not.
5. A mitigation taxonomy comparing server-side aggregation, application-side chunk limits (such as `hyper's BodyExt::limit_chunks(N)` [2]), frontend buffering, and transport-layer rate limiting.
6. A reproducer harness (per-server Dockerfile + standardized probe + JSON results, schema-validated) published with the paper.

Roadmap. Section 2 provides background on the chunked encoding and on the dispatcher architectures we measure. Section 4 describes the measurement protocol. Section 5 treats each server individually. Section 6 compares them. Sections 7–9 analyze root cause, threat model, and mitigations.

2 Background

2.1 HTTP/1.1 chunked transfer encoding

RFC 9112 §7.1 [1] specifies chunked transfer encoding as a sequence of length-prefixed chunks. Each chunk has the form `<hex-size>;<extensions>\r\n<data>\r\n` where `<extensions>` is optional and a chunk of size 0 terminates the message. A trailer section MAY follow the zero chunk before the final `\r\n`.

The encoding was introduced in RFC 2616 (1999) to support sending content whose length is not known when the request line is sent. It is used by streaming clients (cURL with `-data-binary @-`, every modern browser forwarding a streamed body, every LLM SDK streaming a prompt) and by the upload paths of essentially every public-facing web framework.




Figure 2: Wall time vs chunk count N under Mode A (bridge-coalesced). Log-log axes; reference slopes at 1 (linear) and 2 (quadratic). The Node trio (`http.Server`, `express`, `fastify`) are the three data series that approach the quadratic reference; `bandit` stays on the linear slope at a high constant.

The spec defines the *wire* format but does *not* constrain how a recipient should schedule decoded body bytes into application code. A recipient may legally:

- Deliver each decoded chunk’s bytes as a separate application-receive event (*per-chunk*; what most servers in our sample do).
- Coalesce all chunk bytes available in the current TCP `recv()` buffer into one application-receive event (*per-recv*; what `System.IO.Pipelines/Kestrel` does).
- Fully buffer the entire body up to a configured cap before invoking the handler (*fully-buffered*; what Puma’s Tempfile path does at the handler boundary, but the *parser* still loops per chunk).

This choice is the load-bearing decision behind everything in this paper.

2.2 HTTP server concurrency models

The servers we measure span the major concurrency paradigms:

Event loop, single-thread. One OS thread, an event loop (`epoll/kqueue/io_uring`), and cooperative tasks. Examples in our sample: `uvicorn` on Python’s `asyncio`, `hyper/axum/actix-web` on Tokio, `node http.Server` on `libuv`, `Vert.x` on Netty, `nginx`, `HAProxy`.

Thread per connection. A worker thread blocks on each socket’s body read. Examples: `Apache httpd mpm_event` (per-conn worker), Tomcat (Servlet model), `gunicorn sync workers`, `waitress`, `unicorn`.

N:M scheduler. Many lightweight tasks (goroutines, fibers, BEAM processes) are multiplexed onto fewer OS threads by the runtime. Examples: Go net/http (goroutines), Cowboy (BEAM processes), Falcon on Ruby (async fibers).

Actor model. Message-passing isolation per request. Example: `actix-web` (Actix actor system on top of Tokio).

The per-chunk amplification class does *not* discriminate among these paradigms. We find it in all four, with broadly comparable cost constants once normalized for runtime overhead.

2.3 TCP coalescing and the measurement two-mode dichotomy

A naive measurement of "how slow is the chunked path" depends critically on whether wire chunks arrive at the server kernel *one per packet* or *batched many per packet*. Both scenarios are realistic:

- In a kubernetes pod-to-pod call across a docker bridge, the kernel will coalesce dozens to hundreds of small TCP segments into one `recv()` of ~4 KiB.
- In a slow-drip / slow-loris-style attack, or behind a CDN that does not aggregate, each chunk can land in its own `recv()`.

The two regimes can differ by orders of magnitude on the same server. We therefore measure both, denoting them *Mode A* (bridge-coalesced) and *Mode B* (paced, each chunk in its own `recv()`); see Section 4.

2.4 Slow-loris lineage and CWE-407

The attack pattern of "send small request fragments slowly to keep a server busy" goes back to Apache Killer / slow-loris and the RUDY family. The class is catalogued as CWE-407 (Inefficient Algorithmic Complexity) [3]; recent ecosystem-relevant CVEs include CVE-2025-67725 (Tornado quadratic header coalescing) [4], CVE-2025-14550 (Django ASGI repeated-header concatenation) [5], CVE-2026-7790 (cowlib chunk-size hex amplification) [6], CVE-2023-39326 (Go net/http chunk-extension cost cap) [7], and the older CVE-2014-0075 (Tomcat chunk-size integer overflow) [8].

Each of these CVEs targets a *specific* cost in the chunked-TE pipeline (per-byte hex parsing, per-header coalescing, integer-overflow on chunk size). None targets the per-chunk application-dispatch cost we measure here. The bug class we characterize is the same family but a previously-unnamed shape; Section 3 positions this lineage against the broader HTTP DoS literature (slow-DoS, algorithmic-complexity, and the modern HTTP/2 control-plane and decompression attacks).

3 Related work

ChunkLoris sits at the intersection of two long lineages: low-bandwidth application-layer denial of service (the slow-DoS family) and algorithmic-complexity / resource-asymmetry attacks. Our contribution is not a new primitive but a *measurement and systematization*: we identify and quantify a resource axis—per-DATA-unit application *dispatch* CPU cost—that the prior literature names only in passing, and we measure it across 27 production servers. We organize the prior work by the resource each attack amplifies, and draw the line to ours in each case.

3.1 Connection-slot exhaustion: the slow-DoS family

The canonical low-and-slow attacks amplify *connection slots*: an attacker holds open many concurrent connections at negligible client cost, exhausting a server's bounded pool. Slowloris [9] (Hansen, 2009) keeps thread-per-connection servers' worker pools occupied by dribbling partial request headers and

never completing them; RUDY (R-U-Dead-Yet) [10] does the same with a slow POST body under a large `Content-Length`; slow-read variants stall on the response side. Cambiaso et al. [11] give the standard taxonomy of these attacks (pending-request, long-response, and multi-layer slow DoS). Section 2.4 situates ChunkLoris in this lineage at the wire level.

The distinction is the bottleneck. The slow-DoS family is *concurrency-bound*: a coalescing front-end or a per-connection header/idle timeout neutralizes it, and once admitted the request costs the server essentially nothing. ChunkLoris is *CPU-bound*: a single admitted, RFC-compliant, fully-framed chunked request burns single-core CPU *in proportion to its chunk count* on the parser/dispatcher boundary, independent of how many connections the attacker holds. Header timeouts and slot caps do not bound it; only per-request work or byte accounting does.

3.2 Algorithmic-complexity and asymmetric-resource attacks

Crosby and Wallach [12] introduced low-bandwidth algorithmic-complexity attacks: worst-case inputs that drive a data structure (hash tables, binary trees) into its degenerate $\Theta(n^2)$ regime, converting a small request into large server CPU. This is the conceptual ancestor of the superlinear case we measure (Node’s `http.Server` via `process.nextTick`). MITRE catalogues this class as CWE-407 (Inefficient Algorithmic Complexity) [3], the asymmetry itself as CWE-405 (Asymmetric Resource Consumption (Amplification)) [13], and the umbrella as CWE-400 (Uncontrolled Resource Consumption) [14].

ChunkLoris’s *common* (linear) case is not an algorithmic-complexity bug in the Crosby–Wallach sense: the per-chunk cost is $\Theta(N)$ in chunk count with a large constant, not a degenerate worst case of an otherwise sub-quadratic structure. The amplified resource is the *fixed per-dispatch overhead*—a parser callback plus an application-receive event (and, on most event loops, a task/await switch)—multiplied by an attacker-chosen frame count. It is best read as CWE-405 asymmetry whose unit of amplification is the protocol’s smallest framing unit rather than a hash collision. The three Node servers additionally exhibit CWE-407 behavior, which we treat as a distinct, server-specific finding layered on top of the universal linear tax.

3.3 Control-plane and decompression attacks in HTTP/2

The modern HTTP/2 DoS literature amplifies resources other than per-DATA-unit dispatch, which sharpens our novelty claim. *Rapid Reset* (CVE-2023-44487) [15, 16] amplifies *control-plane stream churn*: a client opens and immediately RST_STREAMs streams, decoupling expensive server-side request setup from the 100-concurrent-stream cap and producing record-breaking request rates. The *CONTINUATION Flood* (CVE-2024-27316 and siblings) [17, 18] amplifies *memory and header-decode CPU* via an unbounded stream of CONTINUATION frames that are decoded but never finalized. *MadeYouReset* (CVE-2025-8671) [19, 20] revives the Rapid Reset control-plane axis by inducing *server-initiated resets* through malformed control frames, thereby sidestepping client-reset rate limits. The calif.io “hidden HTTP/2 bomb” [21] composes an HPACK indexed-reference bomb (one wire byte per reference expands to tens-to-thousands of bytes of server allocation) with a flow-control window stall, amplifying *memory via header decompression* and pinning it with zero-window WINDOW_UPDATE drips.

None of these targets per-DATA-frame application dispatch. Rapid Reset and MadeYouReset never send DATA; CONTINUATION Flood and the HPACK bomb live entirely in the header/control plane. Crucially, HTTP/2 flow control bounds *bytes* in flight but places no ceiling on the *number of frames* carrying those bytes—so an attacker can split one flow-control-window’s worth of body into thousands of 1 B DATA frames, each forcing a parser callback and a handler wake. Our Wave 3 results show every HTTP/2 server in our sample inheriting the same per-frame tax (4 μ s to 10 μ s/frame) as HTTP/1.1 chunked-TE; the per-chunk axis is orthogonal to, and unmitigated by, the defenses deployed for the attacks above.

3.4 HTTP parsing and smuggling measurement lineage

Methodologically, our cross-server probing of how independent implementations diverge on an underspecified part of the HTTP grammar follows the request-smuggling measurement tradition, most prominently Kettle’s *HTTP Desync Attacks* [22], which exploits *disagreement* between front-end and back-end parsers over message framing (Content-Length vs. Transfer-Encoding). That line weaponizes *semantic* parser divergence to smuggle requests; we instead measure *performance* divergence on a point the spec leaves open—the granularity at which a conformant parser delivers decoded body bytes to the handler (Section 2.1). We share the comparative, many-implementation methodology but target a CPU cost axis rather than a confusion-based exploit, and every server we measure is fully RFC-compliant.

4 Methodology

4.1 Test harness

Every server is built as a Docker image pinning the exact version under test and exposing port 8000 on a Docker bridge network. The container is constrained to a single CPU (`-cpus=1`). The application handler implements the minimum needed for a meaningful measurement: a GET `/health` returning 200 OK with body `{"ok":true}` and a POST `/upload` that fully drains the request body and returns 200 OK with body `{"len":N}`. No logging middleware, no metrics, no compression, no body-size limit is configured beyond the framework default.

A separate *prober* container, on the same Docker bridge network, reaches the server by its DNS alias and issues the controlled HTTP requests. The prober is a single Python script that opens a raw `asyncio` TCP connection, writes the request headers, then writes the body as chunked-TE chunks per the test mode (see below), and reads the response. Wall time is measured client-side from the moment the prober writes the first byte of the body to the moment the server’s response is received.

4.2 Standardized payload

For each measurement cell we send a body of N bytes, encoded as N chunks of one byte each:

```
POST /upload HTTP/1.1
Host: x
Transfer-Encoding: chunked
Content-Type: application/octet-stream
Connection: close
```

```
1\r\nA\r\n1\r\nA\r\n... (N times) ...0\r\n\r\n
```

Each chunk is `1\r\nA\r\n` (6 wire bytes per data byte), so the bytes on the wire are roughly $6N$ plus a fixed header overhead. We measure at $N \in \{50000, 100000, 250000\}$, with a few servers extended to $N=1000000$ for the log-log scaling charts in Figure 2.

This payload is RFC-compliant. It is a degenerate-but-legal chunked-TE encoding; no protocol-level violation is required.

4.3 Two probe modes

The TCP coalescing behavior on the server depends on how the prober delivers bytes to the kernel. We measure both extremes:

Mode A: bridge-coalesced. The prober writes all chunks back-to-back to its `StreamWriter`; `TCP_NODELAY` is *not* enabled. The kernel coalesces the writes into MSS-sized TCP segments. The server’s `recv()` buffer receives many chunks per syscall return. Empirically, on our Docker bridge this is ~ 8 –16 wire chunks per `data_received` callback.

Mode B: paced 100 μ s. Between every chunk write the prober busy-waits for 100 μ s (a CPU spin, not `asyncio.sleep`; the latter yields to the event loop and is too coarse). `TCP_NODELAY` is enabled. Each chunk leaves the prober as its own TCP segment with $\geq 100 \mu$ s between sends; the server typically returns from each `recv()` with one chunk’s worth of bytes.

Mode A and Mode B bound the realistic spectrum of network conditions: Mode A approximates kubernetes pod-to-pod traffic, Mode B approximates slow-drip attacker pacing or trans-WAN bursts.

4.4 What the per-chunk cost number means

Primary observable. The one quantity we record on every cell is *prober-side wall time* T_{wall} : elapsed seconds from the prober writing the first body byte to receiving the server’s response. This is an end-to-end elapsed-time measurement, not a CPU-time measurement.

Per-chunk cost, defined per mode. For a cell of N one-byte chunks we report a per-chunk cost c in microseconds. Because the two modes load the server differently, T_{wall} contains different components, so we define c per mode:

- **Mode A (bridge-coalesced):** no inter-chunk delay, so T_{wall} is dominated by server work plus scheduling. We report $c_A = T_{\text{wall}} \cdot 10^6 / N$, a wall-derived upper bound on per-chunk server cost (it includes any prober-side and network time, and for the quadratic servers it is the steady-state slope, not a constant).
- **Mode B (paced 100 μ s):** the prober inserts a 100 μ s busy-wait per chunk, so the floor of T_{wall} is $N \cdot 100 \mu$ s whenever the server is faster than the pacing. The per-chunk *server* cost is then the excess over that floor, $c_B = (T_{\text{wall}} - N \cdot 100 \mu\text{s}) \cdot 10^6 / N$. When the server’s own per-chunk cost *exceeds* 100 μ s the pacing is fully overlapped and no longer the floor; for those saturated cells (`cpu` $\approx 100\%$, e.g. `nginx`, `bandit`) we instead report $T_{\text{wall}} \cdot 10^6 / N$, a wall-time upper bound rather than a clean CPU figure.

Wall-derived vs. CPU-derived cells. Where the container is saturated for the full run (`cpu` $\approx 100\%$), wall and CPU time coincide to within measurement error and we treat the figure as CPU. Otherwise wall time over-counts CPU, so the per-chunk number is a wall-time upper bound. Of the 203 cells in our dataset, only 64 carry a server CPU sample (`server_cpu_pct`); the other 139 are wall-derived. We never convert a wall-derived cell into a CPU-second claim, and Section 11.2 records the resulting bias directions.

CPU sampling method. Where present, `server_cpu_pct` is a single `docker stats --no-stream` snapshot taken during the run, and the CPU-seconds we cite are $T_{\text{wall}} \times \text{server_cpu_pct}$. This is a sampled rate times a duration, not an integral of `cpuacct.usage`; under contention `docker stats` can over- or under-report and can exceed 100 % on a single pinned core, so we round CPU figures to the tens-of-percent level.

4.5 Profile collection

For each server we sample its CPU profile during the worst Mode B cell, using whichever profiler the language ecosystem provides:

- Python: `cProfile` with `sort_stats('cumulative')`, top 30.
- Go: `net/http/pprof` exposed on a side port, sampled via `go tool pprof -top`.
- Rust: `perf record -F 999, converted to perf report --stdio(kernel.perf_event_paranoid blocked profiling on some Rust runs; we fall back to source-level analysis for those).`

- Node: `node -prof`, then `node -prof-process top 100`.
- JVM: `async-profiler` agent attached, `itimer` sampling.
- BEAM: `:eprof` posted to a `/profile_upload` endpoint inside the same process.
- .NET: `dotnet-trace` via a sidecar container sharing the target's PID namespace.

4.6 Reproducibility

Every test harness is checked into the project tree under `wave{1,2}/<ecosystem>/repro/`, including the Dockerfile(s), the application source, the prober Python script, the per-cell run logs, and the saved profile artifacts. The top-level rebuild driver, `scripts/render_all.sh`, regenerates every figure from `data/wave*.json` via the chart-rendering scripts and re-runs the LuaLaTeX build of this paper. Anyone with Docker plus `uv` plus `lualatex` can reproduce every number in this paper from the published harness.

5 Per-framework deep dives

This section treats each ecosystem in turn with one subsection per ecosystem. Within each ecosystem, we present per-server measurements, the source-code location of the chunked-decoder dispatch, the dominant profile hot path under attack, and a brief discussion of why the server is in the verdict class it occupies. Numerical results are auto-extracted from `data/wave*.json`; the full reproducer is at `wave{1,2}/<ecosystem>/repro/`.

A 17-dimension per-framework deep dive (versions, parser provenance, source map, profile, mitigations, regression, etc.) is in progress and will appear in the extended version of this paper.

5.1 Python (Wave 1 + extended Wave 2)

ASGI servers. The Python ASGI ecosystem has five widely-deployed servers. `uvicorn` 0.32.1 with the `httptools` backend is the fastest in our sample (2.1 μ s/chunk); the `h11` backend on the same `uvicorn` is 5 μ s. `daphne` (Twisted-based, used by Django Channels) sits at 2.5 μ s. `hypercorn` `h11` backend is 7.5 μ s. `granian` (Rust core via `PyO3`) is the slowest at 14.4 μ s/chunk; the Rust-to-Python bridge cost per chunk dominates its measurement.

All five emit one ASGI `{type: "http.request"}` receive event per `h11` / `httptools` data event. The relevant code paths are `uvicorn/protocols/http/h11_impl.py:257-261` (`uvicorn` `h11` backend) and `uvicorn/protocols/http/httptools_impl.py:301` (`uvicorn` `httptools` backend). None aggregate.

WSGI / hand-rolled servers (Wave 2). `gunicorn` sync workers (Django/Flask production default) register 2.7 μ s/chunk on the fast Mode A but super-linear scaling under heavier loads; the hot path is `gunicorn/http/body.py:77 parse_chunk_size`. `waitress` 3.0.2 reaches 18.4 μ s/chunk under Mode A with similar super-linear behavior. `tornado` 6.4.2 is comparable. None of the three exhibit a clean BATCHES-CORRECTLY pattern: while `gunicorn`'s `'wsgi.input.read()'` does fully buffer at the handler boundary, the chunked decoder underneath still loops one Python function call per chunk before the buffered IO returns.

5.2 Go

The Go standard library's `net/http` chunked decoder lives at `net/http/internal/chunked.go`; its `chunkedReader.beginChunk` runs one `bufio.Reader.fill` call per chunk header. The hot path shows 73% of cumulative time in `syscall.Read + bufio.fill`: one syscall per chunk. `gin` is indistinguishable because it uses the `stdlib` decoder unchanged. Both servers measure 13.4 μ s/chunk under Mode B.

Go has shipped one adjacent CVE in this family (CVE-2023-39326 [7]: a chunk-extension cost cap), but the raw-chunk-count amplification we measure is explicitly not covered by the existing cap.

5.3 Rust

`hyper`'s internal chunked decoder (`ChunkedState::Body`) yields one `Frame::data(Bytes)` per parsed chunk. `axum` (built on `hyper`) sits at 4.3 $\mu\text{s}/\text{chunk}$. `actix-web` (different concurrency model: Actix actor system on top of Tokio) uses its own `actix-http PayloadDecoder` but the same one-frame-per-chunk pattern; it measures 2.3 $\mu\text{s}/\text{chunk}$.

The Rust ecosystem also provides the closest thing to public prior art on this class: `hyper` issue #4008 [2] was opened as an RFC for built-in chunked request limits and closed `not_planned` on 2026-01-12 with the labelled state `S-waiting-on-author`. The maintainer's closing comment, in substance, was that the maintainer was personally skeptical of the described problem ("if as a developer my server does heavy work, I need to design it properly") and encouraged the reporter to maintain a `limit-helpers` crate as a separate library. The proposed `BodyExt::limit_chunks(N)` helper was not shipped in `http_body_util`. We characterize this as a feature-scope decision by a maintainer skeptical of the threat framing; we do not characterize it as a deliberate security position.

5.4 Node.js (the three QUADRATIC servers)

Node's built-in `http.Server` measures 1128 $\mu\text{s}/\text{chunk}$ at $N=250,000$ (Mode A, measured), with a measured scaling exponent of $\sim 1.7\text{--}2.0$ in the $N \in \{50,000, 100,000, 250,000\}$ range. `express` 5.x and `fastify` 5.x extrapolate from smaller cells (measured $N=50,000$ and $N=100,000$) to ~ 324 s and ~ 339 s at $N=250,000$ respectively¹. The framework layer adds essentially nothing because all three use the same underlying `http.Server + llhttp` substrate.

The profile hot path is: `llhttp on_body (C) -> parserOnBody (node:_http_common:128) -> readableAddChunkPushByteMode (node:internal/streams/readable:463) -> addChunk (node:internal/streams/readable:520) -> maybeReadMore (node:internal/streams/readable:857) -> process.nextTick (node:internal/process/task_queues:140)`.

The `nextTick` call per chunk is the quadratic mechanism. Each scheduled callback drains a prefix of the microtask queue; the queue grows with the number of pushed chunks, so the drain work per scheduled callback grows linearly, giving $\Theta(N^2)$ total work.

5.5 JVM

`Vert.x` (Netty-based) measures 7.2 $\mu\text{s}/\text{chunk}$ under Mode B; the hot path is `HttpObjectDecoder.decode -> Http1xServerRequest.onData`, one `onData` call per Netty `HttpContent` fragment. Netty's `HttpObjectAggregator` [23] is the built-in primitive for batching, but `Vert.x` does not install it on the default codec pipeline because doing so would break the streaming `request().handler()` contract.

`Spring Boot` 3.3.5 with the embedded `Tomcat` 10.1 measures 21 $\mu\text{s}/\text{chunk}$ – the worst linear cell in our sample (excluding `nginx-as-origin`). The extra cost over `Vert.x` is the per-chunk `Tomcat NIO Poller -> worker SynchronizedQueue` thread handoff, visible as `pthread_cond_signal` and `ObjectMonitor::wait` traffic in `async-profiler`.

5.6 Ruby

`Puma` 6.4.3 measures 29 $\mu\text{s}/\text{chunk}$ despite a threaded model with `Tempfile` buffering at the handler boundary; the per-chunk decode loop is pure-Ruby `Puma::Client#decode_chunk (lib/puma/client.rb:526-630)`.

¹Both the `express` and `fastify` $N=250,000$ cells are extrapolated from a quadratic fit and are marked with an asterisk in Table 1; only the underlying $N=50,000$ and $N=100,000$ data points were measured directly. The directly-measured 1.5 MiB headline figure we cite (282 s of single-core wall time for Node `http.Server`) is a measurement at $N=250,000$; the `express/fastify` 324/339 s figures are not, and are excluded from any measured total.

Unicorn 6.1.0 beats Puma at 15 μ s/chunk by virtue of a Ragel-generated C extension (`ext/unicorn_http/unicorn_http`). Falcon 0.49.0 (async fiber model on `async-http`) is the cheapest in the Ruby cohort at 10 μ s/chunk.

A prior Puma advisory, CVE-2024-21647 / GHSA-c2f4-cvqm-65w2 (chunk-extension size cap, fixed in 6.4.2/5.6.8), is structurally similar to Go's CVE-2023-39326: it caps a related cost but does not address raw chunk count.

5.7 BEAM (Erlang / Elixir)

Cowboy 2.15.0 (with `cowlib` 2.16.1) measures 14 μ s/chunk; the cost is the Erlang process message-send per chunk from `cowboy_http` (the connection process) to `cowboy_stream_h` (the per-request handler process). Phoenix 1.7 on top of `plug_cowboy` adds the Plug adapter layer and measures 32 μ s.

Bandit 1.11.1 (pure-Elixir HTTP/1, no Cowboy) is a high-constant *linear* case ($\sim 138 \mu$ s/chunk under Mode B), *not* quadratic: its measured scaling exponent is ~ 1.05 and its per-chunk cost *falls* with N ($155 \rightarrow 148 \rightarrow 138 \mu$ s over $N=50k \rightarrow 100k \rightarrow 250k$). It nonetheless carries the same "per-chunk operation walks accumulated state" shape as the Node trio: `erlang:iolist_size/1` consumes 89.3% of CPU in the `:eprof` sample, because `Bandit.HTTP1.Socket.do_read_chunk!/4` accumulates body fragments into an iolist and re-walks it on every chunk via `iolist_size/1` to update progress. In our N range that walk stays bounded enough to remain linear; we classify Bandit VULNERABLE-PER-CHUNK at the top of that class.

5.8 .NET (the HTTP/1.1 BATCHES-CORRECTLY reference)

Kestrel 9 (ASP.NET Core 9) on .NET 9 is the single HTTP/1.1 server in our sample that defeats the per-chunk class structurally (four WebSocket servers do the same at the frame layer; Section 5.11). `Http1ChunkedEncodingMessageBody.PumpAsync(src/Servers/Kestrel/Core/src/Internal/Http/Http1Ch` in the `release/9.0` branch) drains the entire socket-readable buffer into a `System.IO.Pipelines.Pipe`, loops the chunked decoder through every chunk header present in the buffer, calls `FlushAsync()` once, and only then yields back to the application's `PipeReader.ReadAsync`. The handler observes one coalesced `ReadResult`.

The `System.IO.Pipelines` substrate was introduced specifically to allow this kind of zero-allocation streaming hot path; the per-chunk amplification class is closed by construction. We measured no GC events in `dotnet-trace` during the worst Mode B cell, consistent with Microsoft's zero-alloc-hot-path claim.

The adjacent CVE-2025-55315 (CVSS 9.9, October 2025) is a Kestrel chunk-extension parser smuggling bug (a parser correctness bug, not an amplification bug), and is the only published Kestrel CVE touching chunked-TE in recent memory.

5.9 C servers (nginx, Apache, HAProxy)

As *origin* servers, all three are VULNERABLE-PER-CHUNK: nginx 1.29 (OpenResty distribution; the per-chunk cost is in `ngx_http_parse_chunked_state_machine`) measures 114 μ s/chunk – among the worst per-chunk cells in our sample (second only to the `bandit` iolist outlier at $\sim 138 \mu$ s), and a synthetic figure rather than a representative nginx deployment: the C cost is amplified by the OpenResty Lua-sink wrapper that we used to obtain a sink-handler in nginx's request lifecycle. nginx's real-world role is as the buffering reverse proxy of Section 9, not as an origin app server. Apache `httpd` 2.4.67 (event MPM, `ChunkedInputFilter`) measures 15 μ s/chunk. HAProxy 3.0.23 (with a Lua applet sinking the body) is the fastest C server at 7.6 μ s/chunk.

As *reverse proxies*, the picture inverts: nginx (`proxy_request_buffering` on) and Apache (`mod_proxy_http` with default `ProxyIOBufferSize 8192`) aggregate the entire body before opening the upstream connection, fully mitigating the per-chunk class for the upstream application server. HAProxy does *not* aggregate by default (see Section 9.4).

5.10 HTTP/2 (Wave 3)

We extended the survey to HTTP/2 by sending the same N -byte body as N one-byte DATA frames over a single h2 stream. We measured six servers covering five language ecosystems: `hypercorn-h2`, `node-http2.Server`, `Kestrel-h2`, `Vert.x-h2`, `rust-hyper-h2`, and `go-h2c`.

All six are VULNERABLE-PER-FRAME at the DATA-frame level, with per-frame costs clustered in 4–10 $\mu\text{s}/\text{frame}$ – the same band as their HTTP/1 chunked-TE measurements (allowing for the extra h2 framing overhead). Three key observations:

1. **HTTP/2 wire amplification is worse per byte than chunked-TE.** A 1-byte DATA frame is 10 wire bytes (9-byte frame header + 1 byte payload); a 1-byte chunked-TE chunk is 6 wire bytes (`1\r\nA\r\n`). HTTP/2 actually *worsens* the bandwidth asymmetry available to the attacker.
2. **Flow control does not rate-limit frame count.** HTTP/2 `SETTINGS_INITIAL_WINDOW_SIZE` defaults to 65 535 octets and `WINDOW_UPDATE` frames auto-grant in every implementation we tested; the window bounds bytes, not frames.
3. **Kestrel’s HTTP/1 batching does not carry over to HTTP/2.** The `Http1ChunkedEncodingMessageBody.PumpAs` loop has no h2 analog; the HTTP/2 framer feeds the Pipe one DATA frame at a time, and Kestrel-h2 becomes per-frame linear like the rest at 4.9 $\mu\text{s}/\text{frame}$. We confirmed this is mechanical rather than a configurable behavior by reading `src/Servers/Kestrel/Core/src/Internal/Http2/`.

No prior CVE or GHSA covers per-frame body-DATA amplification in HTTP/2; published h2 DoS CVEs (Rapid Reset, CONTINUATION Flood, MadeYouReset) target control-plane and header frames.

5.11 WebSocket (Wave 3)

We sent N WebSocket text frames of one character each and measured server CPU per frame. Six servers: `py-uvicorn-websockets`, `py-uvicorn-wsproto`, `node-ws`, `go-gorilla`, `rust-tungstenite`, `dotnet-kestrel-ws`.

The WebSocket picture is structurally different from HTTP/1 and HTTP/2: **4 of 6 WebSocket servers BATCH-CORRECTLY** by default. Each frame parser loops every in-buffer frame before yielding to the application handler:

- `node-ws` (ws 8.18.0): 0.49 $\mu\text{s}/\text{frame}$ Mode A
- `gorilla/websocket` 1.5.3: 0.13 $\mu\text{s}/\text{frame}$
- `tokio-tungstenite` 0.24: 0.09 $\mu\text{s}/\text{frame}$
- ASP.NET Core WebSockets (Kestrel): 0.66 $\mu\text{s}/\text{frame}$ – Pipelines pattern reused

The two outliers are the Python ASGI WebSocket implementations: `uvicorn + websockets` at 4.74 $\mu\text{s}/\text{frame}$ and `uvicorn + wsproto` at 9.49–14.35 $\mu\text{s}/\text{frame}$. Both emit one ASGI `websocket.receive` event per parsed frame, with the coroutine wake-up dominating the cost – the same shape as their HTTP/1 chunked-TE behavior.

The WebSocket result is paper-worthy for two reasons: (a) it shows the bug class is *not* universal – when the frame parser is designed to drain before dispatching, the cost disappears; (b) it identifies Python ASGI WebSocket as an ecosystem-specific outlier in the `websockets / wsproto` stack.

Figure 3 summarizes the per-protocol ranking across the servers we measured in multiple protocols.




Figure 3: Per-protocol per-chunk-or-frame cost for servers measured in at least two of (HTTP/1 chunked-TE, HTTP/2 DATA, WebSocket text frame). Log-scale y. Kestrel batches in HTTP/1 but not HTTP/2; node-ws, gorilla, tungstenite, and Kestrel-WS batch in WebSocket but their HTTP/1 counterparts do not. The batching behavior is per-protocol-implementation, not per-server-author.

6 Cross-framework synthesis

6.1 Master comparison matrix

Table 1 summarizes the per-chunk cost, scaling exponent, and overall verdict for every server measured. Servers are sorted by verdict (BATCHES-CORRECTLY first, then VULNERABLE-PER-CHUNK in ascending per-chunk cost, then QUADRATIC). The verdict column uses the four-level taxonomy from Section 4.3.

6.2 Three structural classes

Across the 27 HTTP/1.1 servers and 7 ecosystems we observe three structural classes:

Class 1: BATCHES-CORRECTLY (1 HTTP/1.1 server). Kestrel on ASP.NET Core. Its `Http1ChunkedEncodingMes` pumps the socket-readable buffer through the chunked-decoder loop entirely inside one synchronous `Read()` invocation, calling `FlushAsync()` only once before yielding back to the application’s `PipeReader.ReadAsync`. The handler observes one `ReadResult` containing the concatenated bytes of every chunk present in the buffer; per-chunk decoder cost is paid in managed code inside the pump and never crosses the `async/await` boundary. `System.IO.Pipelines` is the substrate that makes this efficient.

Class 2: VULNERABLE-PER-CHUNK linear (23 servers). The chunked decoder loops one application callback per chunk (ASGI `await receive()`, Go’s `Body.Read` returning a single chunk, Rust hyper’s `ChunkedState::Body` yielding one `Frame::data(Bytes)`, Node’s `parserOnBody` pushing into a `Readable` stream, BEAM’s `cowboy_stream_h:data/4` message, Ruby’s `decode_chunk`, JVM’s `Http1xServerRequest.onData`). Per-chunk cost is dominated by the dispatch primitive in the runtime

Table 1: Master comparison matrix. Per-chunk cost is the server-CPU cost measured under Mode B (paced 100 μ s, $N=250,000$, 1 vCPU). Auto-generated from data/wave*.json via scripts/render_matrix.py.

#	server	version	ecosystem
1	Kestrel	9.0 (Microsoft.AspNetCore 9.0; bundled in mcr.microsoft.com/dotnet/aspnet:9.0)	dotnet
2	uvicorn (httptools)	0.32.1	python
3	actix-web	4.x	rust
4	daphne	4.x	python
5	axum (on hyper)	0.7 / hyper 1.x	rust
6	uvicorn (h11)	0.32.1	python
7	Vert.x	4.5.10	jvm
8	HAProxy (with Lua applet for body sink)	3.0.23	c
9	hypercorn	0.17.3	python
10	falcon (on async-http / protocol-http1)	0.49.0	ruby
11	net/http (stdlib)	go-1.23.12	go
12	gin	1.10.1	go
13	cowboy	2.15.0 (cowlib 2.16.1, ranch 2.2.0)	beam
14	unicorn	6.1.0	ruby
15	Apache httpd (event MPM + mod_lua)	2.4.67	c
16	granian	1.x	python
17	Spring Boot / Tomcat	spring-boot-3.3.5 / tomcat-10.1	jvm
18	puma	6.4.3	ruby
19	phoenix	1.7 on plug_cowboy 2.7 / cowboy 2.15.0	beam
20	tornado	6.4.2	python
21	gunicorn	23.0.0	python
22	waitress	3.0.2	python
23	nginx (openresty distribution)	1.29.2.3 / openresty:alpine	c
24	bandit	1.11.1	beam
25	node http.Server (built-in)	node-22.22.3	node
26	express	5.2.1	node
27	fastify	5.x	node

rather than the parser itself: ASGI receive, goroutine wake, future-poll, `process.nextTick`, mailbox enqueue, fiber switch, or thread context switch.

The cost-per-chunk spread within this class is $\sim 65\times$, from 2.1 μ s/chunk (uvicorn + httptools) to 138 μ s/chunk (the `bandit` iolist outlier), with the synthetic `nginx-as-origin` cell at 114 μ s/chunk just below it. Compiled languages (Rust, C, Go) cluster at the cheap end; managed-runtime servers running their parser in interpreted code (Puma, gunicorn, waitress, tornado) cluster at the expensive end, and `bandit`'s iolist walk (Section 7.2) is the single highest-constant linear case.

Class 3: QUADRATIC (3 servers). A single superlinear mechanism appears in our sample, exhibited by the three Node `http.Server`-based servers, which all share one substrate:

- *Node.js Readable substrate.* Each parsed chunk causes `llhttp on_body -> parserOnBody -> socket.push(buf) -> addChunk -> maybeReadMore -> process.nextTick`. The `nextTick` per chunk grows the microtask queue; drain cost is superlinear in queue depth. The built-in `http.Server` reaches a measured exponent of $\sim 1.7\text{--}2.0$; `express` and `fastify` inherit the same substrate and the same shape.

`Bandit` carries a related "per-chunk operation walks accumulated state" pattern (an `erlang:iolist_size/1` walk per chunk), but in our N range it does *not* reach the quadratic regime: its measured exponent is ~ 1.05 and its per-chunk cost *falls* with N ($155 \rightarrow 148 \rightarrow 138$ μ s/chunk under Mode B). We therefore classify it VULNERABLE-PER-CHUNK at a high linear constant (~ 138 μ s/chunk), not QUADRATIC.

6.3 Headline ranking

Figure 1 ranks all 27 HTTP/1.1 servers by per-chunk cost on a log-scale x-axis. The most-deployed servers in each ecosystem (**express** in Node, **Spring Boot** on Tomcat in JVM, **Puma** in Ruby, **gunicorn** in Python, **nginx** as origin in C) sit at or near the worst end of their respective ecosystems. This is the dominant pattern we observe across the matrix: *the most-deployed server in each ecosystem is rarely the fastest at this workload*.

Figure 2 plots wall time as a function of N on log-log axes; the three QUADRATIC servers (the Node trio) separate cleanly from the linear pack. Figure 6 contrasts Mode A versus Mode B at $N=250,000$ and shows that for several servers Mode A is two to three orders of magnitude faster than Mode B, hiding the bug from any test methodology that omits paced probes.




Figure 4: Verdict distribution across all 39 surveyed servers (27 HTTP/1.1, 6 HTTP/2, 6 WebSocket). Five batch correctly (Kestrel on HTTP/1.1 plus four WebSocket stacks); 29 are per-chunk vulnerable; 3 are quadratic (the Node `http.Server` trio); 2 WebSocket stacks are per-frame vulnerable. Restricted to the 27 HTTP/1.1 servers the split is 1 / 23 / 3.

7 Root cause analysis

7.1 The parser/dispatcher boundary

Every HTTP server in our sample is structured as a parser feeding a dispatcher. The parser decodes wire bytes (chunked-TE framing, header parsing, body length tracking). The dispatcher delivers decoded body bytes into application code (ASGI handler, Go’s `http.Handler`, Tower service, Node `Readable`, Servlet `getInputStream`, Erlang process mailbox, Pipelines `PipeReader`).

The per-chunk amplification class lives at the boundary between those two abstractions. The parser is required by the wire format to track chunk-level state; this is unavoidable and roughly constant-cost per chunk in our sample (10s of nanoseconds to a few microseconds in C, less than 10 μ s in pure managed code). The dispatcher’s cost *per delivery event* varies wildly with the runtime: a Go goroutine wake costs a few microseconds; an `asyncio await receive` costs perhaps 5 μ s; a Node `nextTick` drains the microtask queue (superlinear in queue depth); a JVM thread context switch costs $\sim 15 \mu$ s.

Whether the dispatcher cost is paid *per chunk* or *per recv batch* is the choice that determines per-chunk amplification. None of the servers in our sample except Kestrel pay it per `recv-batch` by default.

Figure 7 sketches the three structural classes as data-flow diagrams from kernel `recv()` to application handler.




Figure 5: Per-chunk cost grouped by ecosystem (sorted by within-ecosystem ascending cost). Hatched bars are QUADRATIC; green-edged bars are BATCHES-CORRECTLY. Every ecosystem with more than one entry has at least one VULNERABLE-PER-CHUNK server; only .NET, with one entry (Kestrel), is free of the class in our sample.

7.2 Three cost mechanisms

We observe three distinct cost mechanisms across the 23 VULNERABLE-PER-CHUNK + 3 QUADRATIC servers.

Mechanism A: dispatcher-bound (most servers). The parser is fast; the cost is the per-chunk dispatch event. Representative example: uvicorn + h11. Profile shows ~ 5 μ s/chunk, of which the h11 state-machine step is ~ 1 μ s and the `message_event.set() -> await receive()` round-trip plus the ensuing async task switch is ~ 4 μ s.

Most of our sample sits in this regime. Optimizing the parser gives essentially no improvement; optimizing the dispatcher (by batching) gives *all* of the improvement.

Mechanism B: thread-handoff (Tomcat). On top of mechanism A, the thread-per-connection Servlet model forces a Tomcat NIO Poller $\rightarrow$ SynchronizedQueue $\rightarrow$ worker thread handoff per chunk. The `pthread_cond_signal / ObjectMonitor::wait` traffic visible in `async-profiler` adds ~ 14 μ s on top of the underlying `ChunkedInputFilter.parseChunkHeader` cost. Total ~ 21 μ s/chunk – the worst linear cell among the conventional servers at $N=250,000$, once the two high-constant outliers (the `bandit` iolist walk at ~ 138 μ s and the synthetic `nginx-as-origin` cell at ~ 114 μ s) are set aside.

Mechanism C: superlinear accumulator (Node). Some servers carry per-chunk *state* that itself grows with the number of chunks already received. When the per-chunk operation has to inspect that state,




Figure 6: Mode A (bridge-coalesced) vs Mode B (paced 100 μs) at $N=250,000$. For many servers Mode A is two or three orders of magnitude faster than Mode B, hiding the bug under default lab-style network conditions.

the total work can become quadratic. In our sample this manifests in the three Node `http.Server`-based servers: the `maybeReadMore` $\rightarrow$ `process.nextTick` chain schedules a callback per chunk; each scheduled callback drains some prefix of the microtask queue; the queue grows with the number of pushed chunks, the drain work per scheduled callback grows similarly, and total work is $\Theta(N^2)$ (measured exponent ~ 1.7 – 2.0).

Bandit exhibits a structurally similar "per-chunk operation walks accumulated state" pattern – `do_read_chunk!/4` calls `erlang:iolist_size/1` on the accumulating iolist on every chunk (89.3% of CPU in the `:eprof` sample) – but in our N range the cost does *not* grow quadratically: the measured exponent is ~ 1.05 and per-chunk cost *falls* with N ($155 \rightarrow 138$ μs/chunk under Mode B), indicating Bandit chunks the accumulator in a way that keeps the per-chunk walk bounded over this range. We therefore treat Bandit as a high-constant linear case, not a quadratic one. (Adjacent prior art: CVE-2026-7790 [6] in `cowlib`'s chunk-size hex parser is a genuine quadratic in the same library family.)

7.3 Why Kestrel avoids both mechanisms

Kestrel's `Http1ChunkedEncodingMessageBody.PumpAsync` drains the entire socket-readable buffer into a `System.IO.Pipelines` Pipe in a single synchronous loop. The loop transitions through `ParseChunkedPrefix` and `ReadChunkedData` states for every chunk in the buffer, writes the data bytes into the Pipe (chunk metadata is stripped), and only after the buffer is fully drained does the pump call `FlushAsync()`. The application's `PipeReader.ReadAsync` resumes *once* per pump, with a single




Figure 7: Three structural classes observed across the 27 HTTP/1.1 servers. (a) *per-chunk*: the chunked decoder dispatches one application receive event per chunk; the cost is N task-switches. (b) *batches-correctly* (Kestrel / System.IO.Pipelines): the decoder loop pumps all chunks present in the recv buffer into a Pipe ring buffer and triggers one application receive event per recv batch; the per-chunk cost stays inside the pump and never crosses the `async/await` boundary. (c) *quadratic*: each per-chunk dispatch additionally walks an accumulator whose size grows with N , giving $\Theta(N^2)$ total work. In our sample, the three Node `http.Server`-based servers fit this $\Theta(N^2)$ pattern via the `maybeReadMore`-driven microtask queue.

`ReadResult` containing all decoded chunks’ bytes coalesced.

This is mechanism-A elimination by construction. There is no per-chunk dispatch event because the dispatcher (the `async/await` state machine of the handler) is only ticked once per pump. There is no per-chunk accumulator walk because the Pipe’s ring buffer is amortized constant-cost-per-byte. The result is a per-chunk cost (Section 6.2) dominated entirely by the synchronous parser loop, with no runtime tax on top.

The Pipelines design is publicly documented and predates the security framing of this bug class. Among HTTP/1.1 servers it is, to our knowledge, the only production server in widespread use that adopts this pattern as a default; four WebSocket implementations (Section 5.11) independently arrive at the same drain-before-dispatch design.

7.4 Why batching is spec-compliant

We address a possible counterargument: that delivering one application event per chunk is required by the relevant specifications.

- RFC 9112 [1] defines the wire format. It does not specify any application-API granularity.
- The ASGI specification [24] says a server “may send a single `http.request` event with the entire body” or “may send the body as multiple `http.request` events”. The spec explicitly permits coalescing.
- PEP 3333 (WSGI) [25] defines `environ[‘wsgi.input’]` as a file-like object read by the application; the server is free to back it with any buffer it likes (Tempfile in Puma; BytesIO in waitress).
- Java’s Servlet API exposes `getInputStream()`; chunking is invisible to the handler.
- Node’s `IncomingMessage` extends `Readable`, which is documented as preserving chunk boundaries when called in `Buffer` mode, but the docs do not prohibit coalescing inside `push()`.

In every case the spec permits the batching that Kestrel does. The choice not to batch is implementation-internal, not spec-mandated.

8 Threat model

8.1 Attacker model

We assume an attacker with:

- the ability to open one or more TCP connections to the target server (no authentication, no shared key),
- the ability to send well-formed HTTP/1.1 requests with arbitrary chunked-TE bodies,
- modest egress bandwidth (KB/s, not MB/s; a single residential broadband connection suffices),
- no co-location requirement, no on-path position, no TLS private key.

The attacker does not need to bypass authentication. Per the attack model (Section 8.2), the target endpoint is *any* HTTP path that accepts a request body; for most deployments this includes unauthenticated paths.

8.2 Reachability

Per-chunk amplification is reachable on any HTTP/1.1 endpoint that accepts a request body, irrespective of method. POST, PUT, PATCH, and DELETE-with-body all trigger the parser. The endpoint does *not* have to do anything with the body; even a handler that returns 401 / 403 immediately on auth failure has already paid the per-chunk cost by the time it inspects the body.

This makes the attack reachable on:

- unauthenticated APIs (chat-as-a-service, contact forms, upload endpoints, webhook receivers, OAuth callback URLs),
- authenticated APIs that begin reading the body before the auth check completes (the default order in many ASGI / WSGI / Servlet stacks),
- any endpoint reached through a frontend that does not aggregate (Section 9.4: HAProxy, some CDN configurations).

The wire pattern is RFC-compliant; intrusion detection systems keyed on protocol violations will not fire.

8.3 Cost-per-attacker-byte

Each one-byte chunk on the wire is 6 bytes (1\r\nA\r\n). For $N=250,000$ chunks the attacker sends roughly 1.5 MB of wire bytes plus headers; the server costs are summarized in Table 2.

A residential broadband uplink at 20 Mibit s^{-1} sustains roughly 2.5 MiB s^{-1} , comfortably servicing one or two attacker connections per second. Against express ($N=250,000$, Mode A, *extrapolated* cell: 324 s of single-core *wall* time per request, not a measured CPU integral), one such uplink sustained at one request per second per source IP would, if the server stays CPU-bound on one vCPU, draw on the order of a few hundred core-seconds per wall-second from the target. We flag this as a wall-time-based estimate built on an extrapolated cell: the directly-measured Mode B CPU cost of a single Node http.Server request at $N=250,000$ is far smaller (about 1.5 s), and the gap between that and the saturated-core wall figure is exactly why the "saturating any single-region deployment" claim requires a measured-CPU derivation we leave for future work (Section 11.2, Section 11.3).

Table 2: Per-request attacker bandwidth versus server cost, $N=250,000$ one-byte chunks. The cost column mixes two measurement types, marked accordingly: ^c cells are Mode B server-CPU-seconds ($T_{\text{wall}} \times \text{cpu\%}$ from a `docker stats` snapshot); ^w cells are single-core *wall* seconds (server pinned at one vCPU, so wall is an upper bound on CPU); ^A denotes a Mode A cell, ^e an extrapolated cell). Asymmetry (cost-seconds per attacker-MB-on-wire, higher favors the attacker) is exact only for ^c rows; see Section 4.4.

Server	Wire (MB)	Cost (s)	Asymmetry (s/MB)
Vert.x	1.5	1.8 ^c	1.2
Go net/http	1.5	3.4 ^c	2.3
Spring Boot / Tomcat	1.5	5.3 ^c	3.5
nginx as origin	1.5	29 ^c	19
bandit (linear)	1.5	35 ^c	23
Kestrel 9	1.5	25.9 ^w	—
node http.Server (QUADRATIC)	1.5	282 ^{w,A}	—
express (QUADRATIC)	1.5	324 ^{w,A,e}	—

^c Mode B server CPU-seconds (`docker stats` snapshot $\times$ wall). ^w wall-derived upper bound, not CPU. ^A Mode A cell.

^e extrapolated from a quadratic fit, not measured. `uvicorn` rows are omitted: they have no Mode B $N=250,000$ CPU cell.

8.4 Detection and evasion

Detection. The attack signature on the wire is "request with `Transfer-Encoding: chunked` and many small chunks." Most WAFs do not include a rule for this; most observability stacks log HTTP requests at the application layer (where the chunked-TE detail is lost) rather than at the L4 layer (where per-packet timing would expose it). The attack is functionally invisible to a default observability stack.

A per-connection CPU-accounting layer (cgroup CPU controller, BPF per-cgroup timing) would expose the bug, but is not commonly deployed.

Evasion. A capable attacker can vary chunk size between 1 and a few bytes (keeping the per-chunk cost dominant but breaking signature-based detection on minimum chunk size). The attacker can rotate source IPs to dodge per-IP rate limits, and can interleave the malicious requests with legitimate-looking traffic on the same connection under HTTP keep-alive.

8.5 Combined attacks

The per-chunk amplification combines unfavorably with several adjacent vulnerabilities:

- *slow-loris on headers* (CWE-400 class): an attacker can spend seconds sending request headers slowly before chunked-TE body parsing even begins, holding the connection-handling worker in head-of-line block.
- *repeated-header coalescing* (Tornado CVE-2025-67725, Django ASGI CVE-2025-14550): if the target is also vulnerable to one of those, the same connection can trigger both costs.
- *HTTP/2 multiplexing*: a single connection can carry many concurrent streams, each carrying chunked-TE bodies, all feeding the same per-chunk parser. We have not measured this interaction (Section 11) but it likely amplifies the per-connection cost by the maximum concurrent stream count.

8.6 Distributed amplification

A botnet scales the attack linearly with attacker count. We caution that the per-request Node figures above are single-core *wall* time (and, for `express`, extrapolated), not integrated CPU-seconds, so the following is an order-of-magnitude estimate rather than a measured result. At one request per second per IP against an `express`-like target, a 100-IP botnet (trivial cost on commercial residential proxy services)

would, if the target stays CPU-bound on the affected workers, draw on the order of 100× the per-request core-time – hundreds of core-seconds per wall-second. The measured Mode B CPU cost of a single Node `http.Server` request at $N=250,000$ is about 1.5 s, roughly 200× smaller than the saturated-core wall figure; until a quadratic Node server is re-measured with integrated cgroup CPU at this N (Section 11.2), we state the botnet draw only as an upper-bound estimate and do not claim it categorically saturates a multi-region deployment.

9 Mitigation taxonomy

9.1 Taxonomy

We taxonomize mitigations across four layers, from closest to the parser to closest to the wire. Table 3 summarizes coverage per layer; the prose below treats each layer’s strengths and limits.

9.2 Server-side aggregation

The structural fix is to coalesce consecutive chunk events into a single application delivery, as Kestrel does (Section 7.3). Netty exposes the `HttpObjectAggregator` pipeline handler precisely for this purpose [23]: it accumulates `HttpContent` fragments up to a configured byte limit and fires a single `FullHttpRequest` event downstream. The mechanism exists; what is missing in Vert.x, Spring Boot, and every Netty-derivative we tested is that the aggregator is *not installed by default* because doing so would change the streaming semantics that the framework’s public API contracts guarantee.

In other words: the cost of correctness in this class is breaking the streaming contract by default. Frameworks that have made the opposite choice (Pipelines in Kestrel; full buffering in Puma’s Tempfile handler boundary) closed the bug class at the price of that contract. Frameworks that preserve streaming pay the per-chunk dispatch cost for every chunk forever.

We argue (Section 10) that the right default for new frameworks is to aggregate up to a configurable threshold (e.g., 64 KiB or 100 ms since last flush, whichever first), and expose streaming as an explicit opt-in.

9.3 Application-side chunk limit

hyper closed issue #4008 [2] as `not_planned` on 2026-01-12 with the explicit position that chunked-TE amplification is application responsibility, and pointed users at `http_body_util:Limited` (for byte caps) and the proposed `BodyExt:limit_chunks(N)` helper (for chunk-count caps; not yet shipped in `http_body_util` [2]). No other server in our sample exposes an equivalent chunk-count cap. Application code in any other ecosystem that wants to defend against this attack must write its own middleware that counts chunks during body read.

This is functional but not pleasant: every application has to reinvent the cap, and many will not. The hyper maintainer position is internally consistent (server-side aggregation breaks streaming; chunk caps are an app-level concern) but does not scale across an ecosystem where most application code is unaware of the issue.

9.4 Frontend buffering

The most widely-given advice ("deploy behind nginx") turns out to be valid for nginx and Apache, not for HAProxy.

nginx. `proxy_request_buffering` on is the default. With it enabled, nginx reads the entire decoded request body into `client_body_buffer_size` (default 8–16 KB; spills to disk via `client_body_temp_path` beyond that). Only after the body is fully received does nginx open the upstream connection and forward

the body in a single `Content-Length` request. The upstream sees one well-formed request, not a chunked stream. Per-chunk decoder cost is paid by nginx, not by the upstream.

We verified this empirically: an instrumented Python upstream behind nginx 1.27-alpine (default `proxy_request_buffering` on) received an attacker chunked-TE body of $N=50,000$ one-byte chunks as a single 50 165-byte `recv()` carrying a `Content-Length: 500000` header and no `Transfer-Encoding` header. The same probe against the same upstream with `proxy_request_buffering` off was forwarded as the original chunked stream (five separate `recv()`s, no `Content-Length`, `Transfer-Encoding: chunked` preserved). The full reproducer is at `wave2/c/nginx-as-proxy/` and the captured upstream logs at `wave2/c/nginx-as-proxy/logs/scenario-{A,B}-*.log`.

For every Python, Node, Ruby, Java, or .NET application deployed behind a default nginx, this is full mitigation of the per-chunk amplification class for the upstream. Nginx itself still pays the parsing cost (Section 6.1 reports a synthetic nginx-as-origin cost of $\sim 114 \mu\text{s}/\text{chunk}$ under an OpenResty Lua sink, among the highest per-chunk cells we measured) but nginx is a hardened C server with well-tuned worker concurrency, and as a buffering proxy it never delivers the body per-chunk to the upstream at all, so the impact is bounded.

Apache httpd. `mod_proxy_http` behaves the same way: `ProxyIOBufferSize` (default 8192) bounds per-request buffering. Upstream sees a coalesced body.

HAProxy. HAProxy streams by default. `http-buffer-request` can be enabled but only waits for one buffer's worth of bytes (`tune.bufsize`, default 16 KiB), not a full per-chunk accumulator. Upstream servers behind HAProxy receive chunked-TE requests directly and inherit the per-chunk cost. We have not seen this clearly stated anywhere in the HAProxy documentation; operators relying on "HAProxy is in front of my app" as a mitigation are mistaken.

Cloud frontends. We have not measured Cloudflare, Fastly, AWS ALB, GCP HTTP(S) LB, or Akamai. Each of these CDNs and load balancers has its own defaults for chunked-body buffering. We flag this as future work (Section 11). Operators should empirically verify whether their frontend buffers chunks before assuming a mitigation is in place.

9.5 Transport-layer rate limiting

A coarser fallback: rate-limit incoming TCP packets per source IP or per connection. This is universally available (`iptables`, `nftables`, eBPF, dedicated DDoS appliances) but expensive to operate correctly:

- Slow legitimate uploads (cellular networks, satellite, constrained IoT) will be caught by aggressive packet-rate caps.
- Per-IP rate limiting fails open under distributed attacks (botnets, residential proxies).
- Per-connection caps catch low-bandwidth slow-drip but not high-bandwidth bursts.

We recommend transport-layer rate limiting only as a defense in depth, not as the primary mitigation.

9.6 Summary recommendation per deployment tier

Direct-exposed application server. If you cannot deploy a frontend, you must apply Server-side aggregation (Kestrel only) or Application-side chunk limits (hyper, or custom middleware everywhere else).

Behind nginx or Apache (default config). Mitigated for the upstream. Verify `proxy_request_buffering` on (nginx) or `ProxyIOBufferSize` default (Apache).

Behind HAProxy. Not mitigated; you must apply application-side caps or enable a body-buffering proxy in front of HAProxy.

Behind a CDN. Verify with the CDN provider that chunked-TE bodies are aggregated before being forwarded to your origin. Most providers do not document this behavior explicitly.

10 Recommendations

We organize recommendations by audience.

10.1 For server maintainers

1. Adopt the Kestrel pump pattern by default. Concretely: after `recv()` returns, loop the chunked decoder through every chunk available in the read buffer, then deliver *one* application-receive event with the concatenated data bytes. Cap the per-event delivery by either total bytes (a configurable limit, default 64 KiB) or time-since-last-flush (configurable, default 100 ms). Whichever fires first triggers the dispatch.
2. Expose an explicit `streaming` mode for use cases that genuinely need byte-by-byte delivery (video uploads, long-poll-style hybrids). Make this opt-in.
3. Add a per-request chunk-count cap with a sensible default (e.g., 4096 chunks per request) and a configuration flag to raise it. This guards against the class even when aggregation is disabled.
4. Publish a security-property summary in the server's deployment documentation: what the chunked-TE handling guarantees, what is on the operator, what is on the application.

10.2 For application developers

1. Deploy behind `nginx` or `Apache httpd` with default `proxy_request_buffering on/ProxyIOBufferSize`. This single configuration choice mitigates the class for the upstream regardless of language ecosystem.
2. If the deployment cannot use such a frontend (single-node HAProxy, direct-exposed app server, kubernetes ingress without body buffering), write or import a middleware that counts chunks per request and short-circuits with HTTP 413 or 408 beyond a threshold.
3. Set a body-byte cap (`Content-Length` ceiling, framework body-size limit) sized to your application's actual needs.
4. Monitor per-request CPU. Alert on outliers exceeding 2σ of normal. The amplification produces clear per-request CPU spikes that show up in any reasonable request-level CPU-accounting metric.

10.3 For operators

1. Audit the deployment: which servers are direct-exposed, which are behind `nginx` / `Apache` / `HAProxy` / a CDN. Categorize by Section 9.6 mitigation tier.
2. For HAProxy-fronted services and direct-exposed services, enable per-IP packet-rate limiting at the L4 layer (`nftables`, `BPF`). Set thresholds that catch slow-drip without blocking legitimate slow uploaders.
3. For DDoS-protection vendors (Cloudflare, Akamai, Fastly): request explicit confirmation in writing of whether chunked-TE bodies are aggregated before forwarding to origin. The public documentation is silent on this for most providers.

4. Set up an alert on per-process CPU breaching some threshold for sustained time (e.g., one core saturated for ≥ 30 seconds attributable to one connection). The attack pattern is easy to spot in a process tree but invisible in default access logs.

10.4 For protocol designers

1. Update RFC 9112 or its successor with non-normative guidance on per-chunk delivery granularity, explicitly blessing aggregation as spec-compliant. The current spec's silence has resulted in fifteen years of independent implementations all making the wrong default choice.
2. Standardize a chunk-count rate-limit negotiation primitive at the protocol level (analogous to HTTP/2's `SETTINGS_MAX_HEADER_LIST_SIZE`). A client could advertise "I will send at most N chunks", and the server could refuse connections that exceed it.
3. For HTTP/3 over QUIC, consider whether stream-level pacing requirements can be tightened to make the same attack class structurally impossible (QUIC streams have receiver-side flow control but no chunk-count cap).

11 Limitations and future work

11.1 Limitations

Network conditions. Mode A (bridge-coalesced) and Mode B (paced 100 μ s) bound the realistic spectrum but do not cover every deployment. Particularly: cloud LB-to-pod paths sometimes apply per-flow shapers that approximate Mode B even without an explicit attacker pacing strategy. We have not measured those configurations directly.

Single-CPU constraint. Every server runs in a `-cpus=1` container. This is the single-CPU worst case. Multi-CPU servers fan out work across cores, but the per-connection cost is still serial per the affected worker, so multi-CPU does not change the per-connection ratio – it changes how many concurrent attackers are needed to saturate the deployment.

TLS overhead not measured. Every test is plain HTTP. TLS adds AEAD framing on top of TCP; TLS records buffer multiple bytes by their nature, which could incidentally aggregate chunked-TE chunks at the record boundary. We have not measured the magnitude.

HTTP/3 not measured. The Wave 3 set extends the survey to HTTP/2 DATA frames and WebSocket text frames (Sections 5.10–5.11), but HTTP/3 over QUIC streams is not yet measured. HTTP/3's framing is similar to HTTP/2's DATA frames, so we expect the per-frame amplification class to transfer, but QUIC's separate flow-control and stream model could change the constant or the mitigation surface; we have not measured it.

Cold start vs warm path. JVM tests in particular include some JIT warm-up cost in the first few requests. We report the warm-path numbers (request 2– N of each cell); cold-path measurements would change the absolute timings but not the asymptotic shape.

Production traffic distribution. The probability that real-world traffic includes per-chunk amplification attempts is unmeasured. We have not deployed honeypots or instrumented production servers.

11.2 Threats to validity

We report the measurement-construct limitations explicitly; each notes the likely *direction* of bias.

Single run per cell (no confidence intervals). Nearly every cell is $n=1$; we have no within-cell variance and report no confidence intervals. We therefore make no fine-grained ranking claims between servers whose per-chunk costs are within a small factor of each other. The order-of-magnitude separations we rely on (e.g. `nginx` $\sim 100\times$ `uvicorn`; the Node quadratic vs. the linear pack) are robust to plausible run-to-run noise; sub-factor-of-two differences are not.

CPU sampled via `docker stats` snapshot. Where we have CPU at all (64 of 203 cells), it is a single `docker stats -no-stream` reading times the run duration, not an integral of `cpuacct.usage/cpu.stat`. This is noisier than integrating cgroup counters, can exceed 100 % on a single pinned core, and degrades under host contention. We treat CPU figures as accurate to roughly the tens-of-percent level. A re-run should integrate `cpu.stat usage_usec` deltas instead.

Wall time used where CPU is absent. For the 139 CPU-less cells we report wall-derived per-chunk cost (Section 4.4). On a saturated single-vCPU container wall $\approx$ CPU; otherwise wall *over-counts* CPU by absorbing idle, network, and prober time. Direction of bias: wall $\geq$ CPU, so wall-derived numbers are upper bounds on server CPU. The quadratic headline figures (Node 282 s) are wall on a fully-pinned, saturated core, and should be read as “single-core wall, $\approx$ CPU because saturated,” not as independently measured CPU integrals.

Single co-located host; prober/SUT contention. Prober and server run on one host. Mode B’s prober busy-waits (100 μ s CPU spin per chunk) on the same machine as the 1-vCPU server-under-test; if they share a core, prober spin steals cycles and *inflates* server wall time. Direction of bias: server wall time biased upward.

CFS quota vs. dedicated core. `-cpus=1` imposes a CFS bandwidth quota (100 ms period), not a dedicated physical core (`-cpuset-cpus`). Under quota a CPU-bound run is throttled at period boundaries, adding wall time that is not server work. Direction of bias: inflates Mode B wall (and wall-derived per-chunk cost) relative to a dedicated core; CPU-percentage figures are less affected.

Version pinning and host metadata gaps. Several rows omit the exact runtime/version, and the published dataset does not populate host CPU/kernel metadata, so cross-host reproducibility is not fully specified. Direction of bias: none on relative results, but absolute microsecond figures are host-specific and not portable.

11.3 Future work

HTTP/3, gRPC, and HTTP/2 stream multiplexing. Wave 3 has measured HTTP/2 multiplexed DATA frames and WebSocket per-frame amplification (Sections 5.10–5.11); every HTTP/2 server inherits the per-frame tax, and 4 of 6 WebSocket servers batch correctly. Remaining for future work are HTTP/3 over QUIC streams and gRPC streaming. A separate, important gap is the *multiplexed* HTTP/2 attack: our harness sends one stream per connection, but a single connection can carry the maximum concurrent stream count, each stream feeding the same per-frame parser. We expect this to multiply the per-connection cost by the stream-concurrency limit and to defeat per-IP connection caps; quantifying it is the most consequential measurement we have not yet run.

Multipart per-part amplification. `multipart/form-data` is structurally similar: each part has a header section delivered to the parser callbacks per chunk. Companion work in `findings/starlette/2026-05/STARLETTE` (retracted as a security finding after patch verification but the underlying parser shape is the same) suggests the class transfers.

Production telemetry. We have no data on what fraction of production HTTP traffic includes chunked-TE bodies with small chunks. A honeypot study or a partnership with a high-traffic CDN would let us estimate the real-world prevalence of the attack pattern.

Frontend buffering measurement at CDNs. Cloudflare, Fastly, Akamai, AWS ALB, and GCP HTTP(S) LB each have proprietary chunked-body handling. None publicly document whether they aggregate before forwarding to origin. A measurement campaign across these would significantly improve the operational recommendation in Section 9.6.

Formal model of batching correctness. We have argued informally (Section 7.4) that batching is spec-compliant under RFC 9112 and ASGI / WSGI / Servlet. A formal model – expressed in TLA+ or a similar specification language – would make the argument bulletproof and would let server maintainers verify their own batching implementations.

Detection signatures. Building an open-source IDS / WAF signature for the attack pattern is straightforward (Section 8.4) but requires balancing false positives against attacker evasion. A reference Snort / Suricata rule set, validated on our reproducer harness, would be a useful publishable artifact.

12 Conclusion

We measured 27 production HTTP/1.1 servers across 7 language ecosystems plus three C servers – and, in Wave 3, a further 6 HTTP/2 and 6 WebSocket servers (39 records total) – for their handling of the RFC-compliant pathological input " N -byte body delivered as N one-byte chunked-TE chunks." Among the HTTP/1.1 servers, one (Microsoft's Kestrel via `System.IO.Pipelines`) structurally defeats the per-chunk amplification class; 23 are linear-but-vulnerable with per-chunk costs spanning a factor of $\sim 65\times$; 3 (Node's built-in `http.Server`, `express`, and `fastify`, which share one substrate) exhibit superlinear $\Theta(N^2)$ scaling that pins a single CPU core for minutes per 1.5 MiB of attacker traffic. (Elixir's `bandit` is a high-constant linear case, $\sim 138\ \mu\text{s}/\text{chunk}$ with measured exponent ~ 1.05 , not quadratic.)

The root cause is a parser/dispatcher boundary design choice, not a spec violation. RFC 9112 fixes the wire format but is silent on application-API granularity; ASGI, WSGI, Servlet, and Node's `Readable` all permit coalescing. The cost of *not* coalescing is a per-chunk dispatch event whose constant varies by runtime: ASGI receive, goroutine wake, future poll, `process.nextTick`, mailbox enqueue, fiber switch, thread context switch. The amplification class is not unique to any single runtime; we observe it in every concurrency model we measure.

The "deploy behind nginx" mitigation widely advised in framework documentation is valid for nginx and Apache httpd (`proxy_request_buffering on` / `ProxyIOBufferSize` default), but is *not* valid for HAProxy, which streams the body upstream by default. Direct-exposed application servers, HAProxy-fronted deployments, and any deployment with proxy buffering disabled inherit the full per-chunk cost.

We argue that opt-in batching primitives modeled on Kestrel's Pipelines pump should be exposed by every event-loop HTTP server, and that the operational advice "deploy behind nginx" deserves clearer documentation of its applicability and its limits.

The reproducer harness, raw measurements, and figure-rendering scripts are published with this paper at <https://github.com/mjbommar/chunkloris-paper>. Anyone with Docker plus `uv` plus `pdflatex` can reproduce every number in this paper from the supplied harness.

References

- [1] R. Fielding, M. Nottingham, and J. Reschke. *HTTP/1.1*. RFC 9112, IETF. 2022. doi: [10.17487/RFC9112](https://doi.org/10.17487/RFC9112). URL: <https://www.rfc-editor.org/rfc/rfc9112>.

- [2] hyper maintainers. *hyper issue #4008: provide built-in chunked request limits and CPU-safe streaming helpers*. Closed as not_planned 2026-01-12. 2026. URL: <https://github.com/hyperium/hyper/issues/4008>.
- [3] MITRE. *CWE-407: Inefficient Algorithmic Complexity*. Common Weakness Enumeration. URL: <https://cwe.mitre.org/data/definitions/407.html>.
- [4] Tornado Web Server Maintainers. *Tornado quadratic header coalescing DoS*. GHSA-c98p-7wgm-6p64 / CVE-2025-67725. 2025. URL: <https://github.com/tornadoweb/tornado/security/advisories/GHSA-c98p-7wgm-6p64>.
- [5] Django Project. *Django ASGI repeated-header concatenation DoS*. CVE-2025-14550. 2026.
- [6] Erlang Ecosystem Foundation CNA. *cowlib chunk-size hex amplification*. CVE-2026-7790. 2026. URL: <https://cna.erlef.org/cves/CVE-2026-7790.html>.
- [7] Go security team. *net/http chunk-extension cost cap*. CVE-2023-39326. 2023. URL: <https://pkg.go.dev/vuln/GO-2023-2382>.
- [8] Apache Software Foundation. *Apache Tomcat chunk-size integer overflow*. CVE-2014-0075. 2014.
- [9] Robert Hansen (RSnake). *Slowloris HTTP DoS*. ha.ckers.org (orig.); archived tool. Low-bandwidth, connection-slot-exhaustion HTTP DoS tool. 2009. URL: <https://github.com/XCHADXFQAQ77X/SLOWLORIS>.
- [10] OWASP. *R-U-Dead-Yet (RUDY): slow HTTP POST denial of service*. OWASP Foundation. Slow-POST variant of the Slowloris connection-exhaustion family. URL: https://owasp.org/www-community/attacks/Slow_HTTP_Attacks.
- [11] Enrico Cambiaso, Gianluca Papaleo, Giovanni Chiola, and Maurizio Aiello. “Slow DoS attacks: definition and categorisation”. In: *International Journal of Trust Management in Computing and Communications* 1.3/4 (2013), pp. 300–319. DOI: [10.1504/IJTMCC.2013.056440](https://doi.org/10.1504/IJTMCC.2013.056440). URL: <https://dblp.org/rec/journals/ijtmcc/CambiasoPCA13.html>.
- [12] Scott A. Crosby and Dan S. Wallach. “Denial of Service via Algorithmic Complexity Attacks”. In: *Proceedings of the 12th USENIX Security Symposium*. USENIX Association, 2003. URL: <https://www.usenix.org/conference/12th-usenix-security-symposium/denial-service-algorithmic-complexity-attacks>.
- [13] MITRE. *CWE-405: Asymmetric Resource Consumption (Amplification)*. Common Weakness Enumeration. URL: <https://cwe.mitre.org/data/definitions/405.html>.
- [14] MITRE. *CWE-400: Uncontrolled Resource Consumption*. Common Weakness Enumeration. URL: <https://cwe.mitre.org/data/definitions/400.html>.
- [15] NIST National Vulnerability Database. *CVE-2023-44487: HTTP/2 Rapid Reset*. CVE-2023-44487. 2023. URL: <https://nvd.nist.gov/vuln/detail/CVE-2023-44487>.
- [16] Lucas Pardue and Julien Desgats. *HTTP/2 Zero-Day vulnerability results in record-breaking DDoS attacks*. Cloudflare Blog. 2023. URL: <https://blog.cloudflare.com/zero-day-rapid-reset-http2-record-breaking-ddos-attack/>.
- [17] NIST National Vulnerability Database. *CVE-2024-27316: Apache HTTP Server HTTP/2 DoS by memory exhaustion on endless CONTINUATION frames*. CVE-2024-27316. 2024. URL: <https://nvd.nist.gov/vuln/detail/CVE-2024-27316>.
- [18] Bartłomiej Nowotarski. *HTTP/2 CONTINUATION Flood*. nowotarski.info. 2024. URL: <https://nowotarski.info/http2-continuation-flood/>.
- [19] NIST National Vulnerability Database. *CVE-2025-8671: HTTP/2 MadeYouReset DoS via control frames*. CVE-2025-8671. 2025. URL: <https://nvd.nist.gov/vuln/detail/CVE-2025-8671>.

- [20] CERT Coordination Center. *VU#767506: HTTP/2 implementations vulnerable to “MadeYouReset” DoS through HTTP/2 control frames*. CERT/CC Vulnerability Note. 2025. URL: <https://kb.cert.org/vuls/id/767506>.
- [21] Calif.io. *Codex Discovered a Hidden HTTP/2 Bomb*. Calif.io Blog. 2025. URL: <https://blog.calif.io/p/codex-discovered-a-hidden-http2-bomb>.
- [22] James Kettle. *HTTP Desync Attacks: Request Smuggling Reborn*. PortSwigger Research (Black Hat USA 2019). 2019. URL: <https://portswigger.net/research/http-desync-attacks-request-smuggling-reborn>.
- [23] Netty Project. *Netty HttpObjectAggregator*. URL: <https://netty.io/4.1/api/io/netty/handler/codec/http/HttpObjectAggregator.html>.
- [24] ASGI working group. *ASGI 3.0 specification*. URL: <https://asgi.readthedocs.io/>.
- [25] P. J. Eby. *PEP 3333: Python Web Server Gateway Interface v1.0.1*. 2010. URL: <https://peps.python.org/pep-3333/>.

Table 3: Mitigation taxonomy. *Default-deployed* means the mitigation is on without any operator action; *available* means an operator can enable it but does not by default; *absent* means no equivalent primitive exists in the ecosystem.

Layer	Mechanism	Default-deployed	Available
Server-side aggregation	coalesce N chunk events into one handler delivery	Kestrel (Pipelines)	HttpObjectAggregator (Netty)
Application-side chunk limit	per-request chunk count	<i>none</i>	BodyExt::limit_chunks(N) (Netty)
Application-side byte limit	total body byte cap (Content-Length or otherwise)	most servers	N/A
Frontend buffering	reverse proxy aggregates body before forwarding	nginx (proxy_request_buffering on), Apache (mod_proxy_http)	N/A
Transport-layer rate limit	per-IP packet/sec cap	N/A (operator opt-in)	iptables, nftables, BPF